



# SQL Basic

Prof. Hyuk-Yoon Kwon

<https://sites.google.com/view/seoultech-bigdata>

Most parts are based on slides used in Stanford  
(<http://web.stanford.edu/class/cs145>)

# Contents

---

## ■ Summary of Previous Lecture

- Introduction to SQL
  - SQL introduction & schema definitions

## ■ Today's lecture

- Basic single-table queries
- Multi-table queries

---

# **1. SQL Introduction & Definitions**

# SQL Motivation

---

## ■ Dark times 5 years ago.

- Are databases dead?



## ■ Now, as before: everyone sells SQL

- Pig, Hive, Impala



## ■ “Not-Yet-SQL?”

# SQL is a...

## ■ Data Definition Language (DDL)

- Define relational *schemata*
- Create/alter/delete tables and their attributes

## ■ Data Manipulation Language (DML)

- Insert/delete/modify tuples in tables
- Query one or more tables

DDL {

Table of baby-name data

name	rank	gender	year
Jacob	1	boy	2009
Isabella	1	girl	2009
Ethan	2	boy	2009
Emma	2	girl	2009
Michael	3	boy	2009

DML {

Field names

One row (4 fields)

2000 rows all told

# Tables in SQL

---

## Product

PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks
SingleTouch	\$149.99	Canon
MultiTouch	\$203.99	Hitachi

A **relation** or **table** is a multiset of tuples having the attributes specified by the schema

# Data Types in SQL

## ■ Atomic types:

- Characters: CHAR(20), VARCHAR(50)
- Numbers: INT, BIGINT, SMALLINT, FLOAT
- Others: MONEY, DATETIME, ...

## ■ Every attribute must have an atomic type

- Hence tables are flat

Value	CHAR (4)	Storage Required	VARCHAR (4)	Storage Required
' '	'   '	4 bytes	' '	1 byte
'ab'	'ab  '	4 bytes	'ab'	3 bytes
'abcd'	'abcd'	4 bytes	'abcd'	5 bytes
'abcdefgh'	'abcd'	4 bytes	'abcd'	5 bytes

Study more: <https://dev.mysql.com/doc/refman/5.7/en/char.html>

# NULL and NOT NULL

- To say “don’t know the value” we use **NULL**
  - NULL has (sometimes painful) semantics, more detail later

Students(sid:string, name:string, gpa: float)

sid	name	gpa
123	Bob	3.9
143	Jim	NULL

*Say, Jim just enrolled in his first class.*

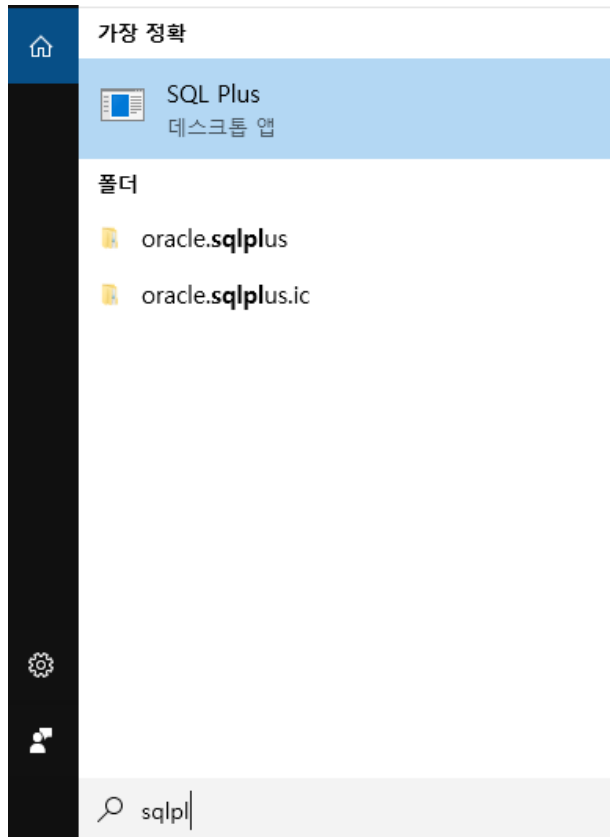
In SQL, we may constrain a column to be NOT NULL, e.g., “name” in this table



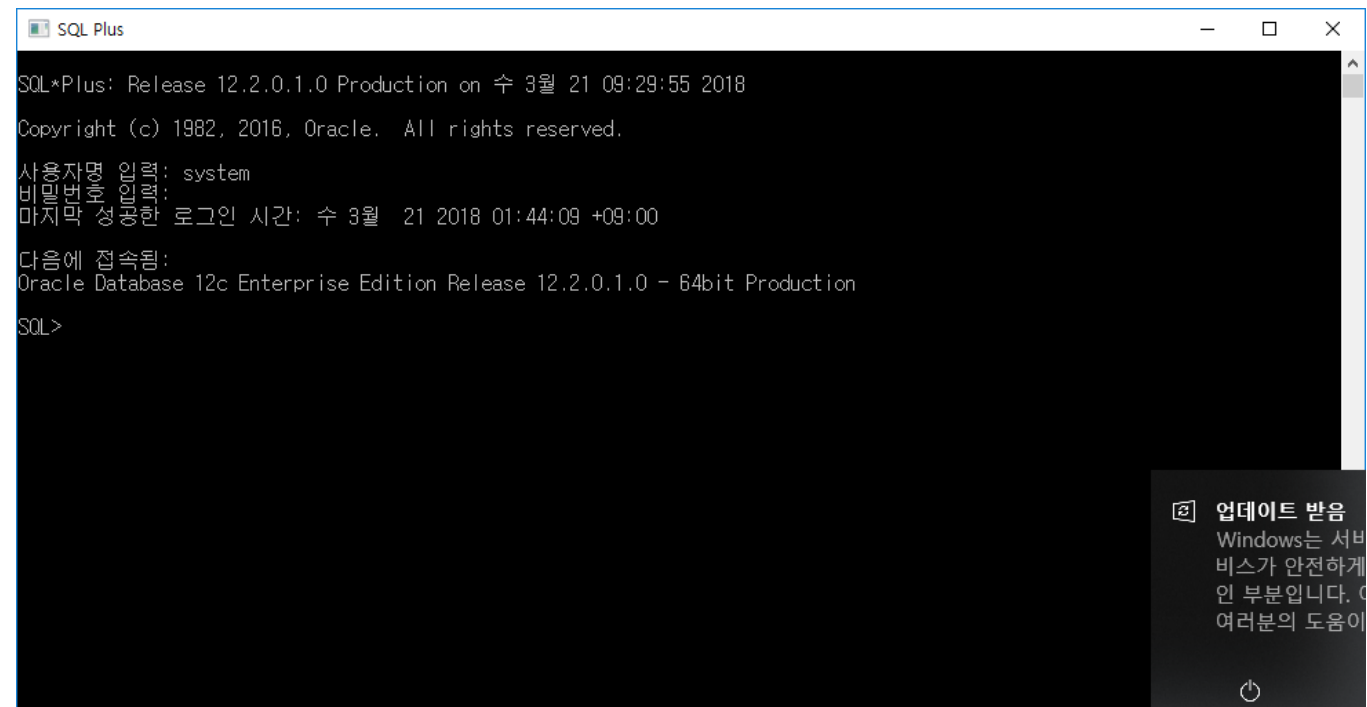
# Oracle Practice #1

## ■ Start Oracle

- Execute SQL Plus



- User Authentication
  - ID: system
  - Password: oraclepractice



# Oracle Practice #1

---

## ■ Build database with real data

1. Download ACDB.sql from e-class
2. Copy ACDB.sql into a specific folder (e.g., c:/work/ACDB.sql)
3. In SQLPlus, execute the following command
  - @c:/work/ACDB.sql
  - If some problems occur, execute the following command, and then execute the command above again
    - alter session set nls\_language="AMERICAN";
4. Check if data are stored correctly
  - select \* from ACDB\_SECTORS;
  - select \* from ACDB\_PACAKGES;
  - select \* from ACDB\_CUSTOMERS;

---

## 2. Single-table queries

# What you will learn about in this section

---

1. The SFW query
2. Other useful operators: LIKE, DISTINCT, ORDER BY

# SQL Query

---

- Basic form (there are many many more bells and whistles)

```
SELECT <attributes>  
FROM   <one or more relations>  
WHERE  <conditions>
```


Call this a **SFW** query.

# Simple SQL Query: Selection

**Selection** is the operation of filtering a relation's tuples on some condition

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM Product  
WHERE Category = 'Gadgets'
```




PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

# Simple SQL Query: Projection

**Projection** is the operation of producing an output table with tuples that have a subset of their prior attributes

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT Pname, Price, Manufacturer
FROM Product
WHERE Category = 'Gadgets'
```



PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks


# Notation

---

Input schema

Product(PName, Price, Category, Manufacturer)

```
SELECT Pname, Price, Manufacturer  
FROM Product  
WHERE Category = 'Gadgets'
```



Output schema

Answer(PName, Price, Manufacturer)



# Database Schema

---

```
CREATE TABLE "ACDB_SECTORS"  
(  
    "SECTOR_ID" NUMBER(8,0),  
    "SECTOR_NAME" VARCHAR2(25 BYTE)  
);
```

```
CREATE TABLE "ACDB_PACKAGES"  
(  
    "PACK_ID" NUMBER(8,0),  
    "SPEED" VARCHAR2(10 BYTE),  
    "MONTHLY_PAYMENT" NUMBER(8,0),  
    "SECTOR_ID" NUMBER(8,0),  
    "STRT_DATE" DATE  
);
```

```
CREATE TABLE "ACDB_CUSTOMERS"  
(  
    "CUSTOMER_ID" NUMBER(8,0),  
    "FIRST_NAME" VARCHAR2(25 BYTE),  
    "LAST_NAME" VARCHAR2(25 BYTE),  
    "CITY" VARCHAR2(45 BYTE),  
    "STATE" VARCHAR2(25 BYTE),  
    "STREET" VARCHAR2(40 BYTE),  
    "MAIN_PHONE_NUM" VARCHAR2(12 BYTE),  
    "SECONDARY_PHONE_NUM" VARCHAR2(12 BYTE),  
    "FAX" VARCHAR2(12 BYTE),  
    "MONTHLY_DISCOUNT" NUMBER(4,2),  
    "PACK_ID" NUMBER(8,0),  
    "BIRTH_DATE" DATE,  
    "JOIN_DATE" DATE  
);
```

# Oracle Practice #2

---

## ■ Given each description, make the corresponding SQL and practice it using ACDB.sql

1. Create a query to display the internet package number, internet speed and monthly payment (*Packages* table).
2. Create a query to display the customer number, first name, last name, primary phone number, secondary phone number and package number (*Customers* table).
3. Display the first name, last name, and package number for all customers whose last name is “King” (*Customers* table).
4. Display the first name, last name, package number and monthly discount for all customers with monthly discount less than 10 (*Customers* table).

# Useful Expressions

a	b
10	20

## ■ Arithmetic operation in SELECT clause

- SELECT a \* 2 FROM table

a
20

## ■ Alias in SELECT clause

- SELECT a as A\_RESULT FROM table

A_RESULT
10

## ■ Concatenate in SELECT clause

- SELECT a || ',' || b as A\_AND\_B FROM table

A_AND_B
10, 20

# Oracle Practice #3

---

## ■ Given each description, make the corresponding SQL and practice it using ACDB.sql

1. Create a query to display the package number, speed, strt\_date (the date when the package became available), monthly payment, and monthly payment \* 12, name the last column "Y\_INCOME" (*Packages* table).
2. Create a query to display the last name concatenated with the first name, separated by space, and main phone number concatenated with secondary phone number, separated by comma and space. Name the column heading FULL\_NAME and CONTACT\_DETAILS respectively. (*Customers* table).
3. Create a query to display the first name, last name, monthly discount and city concatenated with street, separated by space. Name the column headings: FN, LN, DC and FULL\_ADDRESS respectively (*Customers* table).

# A Few Details

---

## ■ SQL commands are case insensitive:

- Same: SELECT, Select, select
- Same: Product, product

## ■ Values are not:

- Different: 'Seattle', 'seattle'

## ■ Use single quotes for constants:

- 'abc' - yes
- "abc" - no

# LIKE: Simple String Pattern Matching

---

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

■ **s LIKE p: pattern matching on strings**


■ **p may contain two special symbols:**

- % = any sequence of characters
- \_ = any single character

# DISTINCT: Eliminating Duplicates

---


```
SELECT DISTINCT Category  
FROM Product
```



Category
Gadgets
Photography
Household

Versus

```
SELECT Category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

# Oracle Practice #4

---

## ■ Given each description, make the corresponding SQL and practice it using ACDB.sql

1. Create a query to display unique cities from the *Customers* table.
2. Create a query to display unique combination of cities and states from *Customers* table.
3. Display the first name and monthly discount for all customers whose first name ends with an *e* (*Customers* table).
4. Display the last name and package number for all customers where the second letter of their last name is *d* (*Customers* table).



# ORDER BY: Sorting the Results

---

```
SELECT PName, Price, Manufacturer  
FROM Product  
WHERE Category='gizmo' AND Price > 50  
ORDER BY Price, PName
```

Ties are broken by  
the second attribute  
on the ORDER BY  
list, etc.

Ordering is  
ascending, unless  
you specify the DESC  
keyword.

# BETWEEN Operation

---

a
20
40

## ■ BETWEEN operation in WHERE clause

- SELECT a FROM table WHERE a BETWEEN 10 AND 30

a
20

# Oracle Practice #5

---

## ■ Given each description, make the corresponding SQL and practice it using ACDB.sql

1. Display the first name, join date, and package number for all customers who don't have the letter *a* in their first name. Order the query in ascending order by package number (*Customers* table).
2. Display the first name, join date, monthly discount, and package number for all customers whose monthly discount is over 28. Order the query in ascending order by monthly discount and package number (*Customers* table)
3. Order the results of the previous problem (#2) in descending order by monthly discount and then in ascending order by package number (*Customers* table)
4. Display the first name, join date, monthly discount where monthly discount is between 28 and 30 (*Customers* table)
5. Display first name and join date where first name is between 'B' and 'C' (*Customers* table)

---

## 3. Multi-table queries

# What you will learn about in this section

---

- Foreign key constraints

- Joins: basics

- Joins: SQL semantics

# Foreign Key constraints

- Suppose we have the following schema:

*Students*(sid: string, name: string, gpa: float)

*Enrolled*(student\_id: string, cid: string, grade: string)

- And we want to impose the following constraint:

- “a student must appear in the Students table to enroll in a class”

**Students**

sid	name	gpa
101	Bob	3.2
123	Mary	3.8

**Enrolled**

student_id	cid	grade
123	564	A
123	537	A+

student\_id alone is not  
a key- what is?

We say that student\_id is a **foreign key** that refers to Students

# Declaring Foreign Keys

---

Students(sid: *string*, name: *string*, gpa: *float*)

Enrolled(student\_id: *string*, cid: *string*, grade: *string*)

```
CREATE TABLE Enrolled(  
    student_id CHAR(20),  
    cid          CHAR(20),  
    grade CHAR(10),  
    PRIMARY KEY (student_id, cid),  
    FOREIGN KEY (student_id) REFERENCES Students(sid)  
)
```

# Foreign Keys and update operations

---

*Students(sid: string, name: string, gpa: float)*

*Enrolled(student\_id: string, cid: string, grade: string)*

## ■ What if we insert a tuple into Enrolled, but no corresponding student?

- INSERT is rejected (foreign keys are constraints)!

## ■ What if we delete a student?

1. Disallow the delete
2. Remove all of the courses for that student
3. *SQL allows a third via NULL (not yet covered)*



# Keys and Foreign Keys

---

## Company

<u>CName</u>	StockPrice	Country
GizmoWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan

What is a foreign key vs. a key here?

## Product

<u>PName</u>	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

# Joins

---

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

*Ex:* Find all products under \$200 manufactured in Japan;  
return their names and prices.

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

# Joins

---

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Ex: Find all products under \$200 manufactured in Japan;  
return their names and prices.

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

A **join** between tables returns all unique combinations of their tuples **which meet some specified join condition**

# Joins

---

Product(PName, Price, Category, Manufacturer)

Company(CName, StockPrice, Country)

Several equivalent ways to write a basic join in SQL:

```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

```
SELECT PName, Price
FROM Product
JOIN Company ON Manufacturer = Cname
              AND Country='Japan'
WHERE Price <= 200
```

A few more later on...


# Joins

Product

PName	Price	Category	Manuf
Gizmo	\$19	Gadgets	GWorks
Powergizmo	\$29	Gadgets	GWorks
SingleTouch	\$149	Photography	Canon
MultiTouch	\$203	Household	Hitachi

Company

Cname	Stock	Country
GWorks	25	USA
Canon	65	Japan
Hitachi	15	Japan



```
SELECT PName, Price
FROM Product, Company
WHERE Manufacturer = CName
      AND Country='Japan'
      AND Price <= 200
```

PName	Price
SingleTouch	\$149.99

# Primary Key and Foreign Key

---

```
ALTER TABLE "ACDB_SECTORS" ADD CONSTRAINT "SECTOR_ID_PK" PRIMARY KEY ("SECTOR_ID");
```

```
ALTER TABLE "ACDB_PACKAGES" ADD CONSTRAINT "PACK_ID_PK" PRIMARY KEY ("PACK_ID");
```

```
ALTER TABLE "ACDB_PACKAGES" ADD CONSTRAINT "SECTOR_ID_FK" FOREIGN KEY ("SECTOR_ID")  
REFERENCES "ACDB_SECTORS" ("SECTOR_ID") ENABLE;
```

```
ALTER TABLE "ACDB_CUSTOMERS" ADD CONSTRAINT "CUSTOMER_ID_PK" PRIMARY KEY ("CUSTOMER_ID");
```

```
ALTER TABLE "ACDB_CUSTOMERS" ADD CONSTRAINT "PACK_ID_FK" FOREIGN KEY ("PACK_ID")  
REFERENCES "ACDB_PACKAGES" ("PACK_ID") ENABLE;
```

# Primary Key and Foreign Key

ACDB\_SECTORS

"SECTOR\_ID"  
"SECTOR\_NAME"

ACDB\_CUSTOMERS

"CUSTOMER\_ID"  
"FIRST\_NAME" "LAST\_NAME"  
"CITY"  
"STATE"  
"STREET" "MAIN\_PHONE\_NUM"  
"SECONDARY\_PHONE\_NUM"  
"FAX"  
"MONTHLY\_DISCOUNT"  
"PACK\_ID"  
"BIRTH\_DATE"  
"JOIN\_DATE"

ACDB\_PACKAGES

"PACK\_ID"  
"SPEED"  
"MONTHLY\_PAYMENT"  
"SECTOR\_ID"  
"STRT\_DATE"

# Oracle Practice #6

---

## ■ Given each description, make the corresponding SQL and practice it using ACDB.sql

1. Write a query to display first name, last name, package number and internet speed for all customers.  
(*Customers* and *Packages*)
2. Display the package number, internet speed, monthly payment and sector name for all packages  
(*Packages* and *Sectors* tables).
3. Display the customer name, package number, internet speed, monthly payment and sector name for all customers (*Customers*, *Packages* and *Sectors* tables).
4. Display the customer name, package number, internet speed, monthly payment and sector name for all customers in the 'Business' sector (*Customers*, *Packages* and *Sectors* tables).


\* You may check the data stored in Sectors table



# Explicit Set Operators: INTERSECT

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
INTERSECT  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cap \{r.A \mid r.A = t.A\}$$



## ■ Constraint for Intersect

```
SELECT Name, BirthDate FROM Employee  
INTERSECT  
SELECT Name, BirthDate FROM Customer
```


```
SELECT Name, BirthDate FROM Employee  
INTERSECT  
SELECT Age, BirthDate, Name FROM Customer
```

# UNION

---

$$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$$

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
UNION  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```



Why aren't there duplicates?


What if we want duplicates?

# UNION ALL

---

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
  UNION ALL  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \cup \{r.A \mid r.A = t.A\}$$



By default:  
SQL uses set semantics

# UNION vs. UNION ALL

1. Employee table data:

```
SELECT * FROM Employee
```

	EmpId	EmpName	EmpCode
1	1	Bhaumik	1609
2	2	Maulik	2568
3	3	Mayur	1254
4	4	Sandip	6578
5	5	Rekansh	7998

2. Customer table data:

```
SELECT * FROM Customer
```

	CustId	CustName	CustCode
1	1	Bhoms	234
2	2	Mayur	656
3	3	Jimit	324
4	4	Sandip	435

3. UNION Example (It removes all duplicate records):

```
SELECT EmpName FROM Employee  
UNION  
SELECT CustName FROM Customer
```

	EmpName
1	Bhaumik
2	Bhoms
3	Jimit
4	Maulik
5	Mayur
6	Rekansh
7	Sandip

4. UNION ALL Example (It just concatenate records, not eliminate duplicates, so it is faster than UNION):

```
SELECT EmpName FROM Employee  
UNION ALL  
SELECT CustName FROM Customer
```


	EmpName
1	Bhaumik
2	Maulik
3	Mayur
4	Sandip
5	Rekansh
6	Bhoms
7	Mayur
8	Jimit
9	Sandip

# MINUS

---

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
MINUS  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```

$$\{r.A \mid r.A = s.A\} \setminus \{r.A \mid r.A = t.A\}$$



# Interesting Result of MINUS!

---

I have 2 tables A and B.

```
SELECT COUNT(*) FROM (SELECT * FROM tableA)
```

returns 389

```
SELECT COUNT(*) FROM (SELECT * FROM tableB)
```

returns 217

```
SELECT COUNT(*) FROM  
(SELECT * FROM tableA  
INTERSECT  
SELECT * FROM tableB)
```

returns 0

```
SELECT COUNT(*) FROM  
(SELECT * FROM tableA  
MINUS  
SELECT * FROM tableB)
```

returns 389

```
SELECT COUNT(*) FROM  
(SELECT * FROM tableB  
MINUS  
SELECT * FROM tableA)
```

# Oracle Practice #7 – Multiset operators

---

## 1. When we have the following two conditions,

- A. IDs of customers who get discount monthly over \$5 (i.e., > 5)
- B. IDs of customers who pay monthly over \$100 (i.e., > 100)
- Find the following results and compare them: 1) A, 2) B, and 3) A INTERSECT B

## 2. When we have the following two conditions, solve the problem

- A. Cities of customers who live in 'California'
- B. Cities of customers who include in 'Business' sector
- Find the following results and compare them: 1) A, 2) B, 3) A UNION B, 4) A UNION ALL B

## 3. When we have the following two conditions, solve the problem

- A. Names of customers who include in 'Private' sector
- B. Names of customers who live in 'Seattle' or 'San Francisco'
- Find the following results and compare them: 1) A, 2) B, and 3) A MINUS B